

Do Pod and Machine Autoscalers work in Harmony?

A performance review of ARO and ROSA

Agenda:

- Context What type of workloads are we talking about?
- Pod Autoscalers

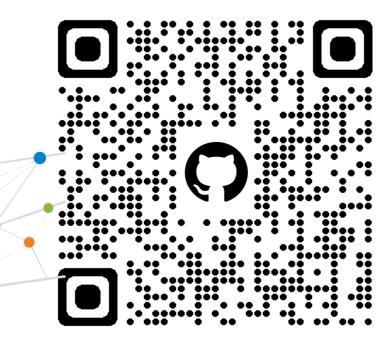


KEDA + Demo

Auto scaling Machines + Demo

ARO

ROSA



What we do & who we are

We enable people to manage the complexity of investment decision making

Ortec Finance is a global technology and solutions provider



Global client base











































































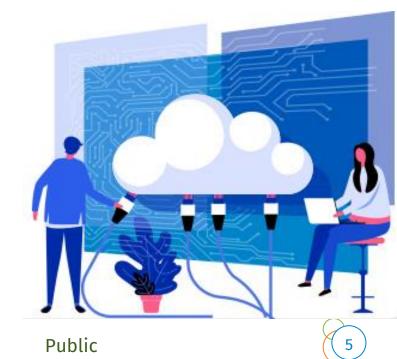






Our Cloud-Native Journey

- We like Managed over DIY:
 - OpenShift is a very mature, battle-tested, enterprise-level Kubernetes toolkit. It's a 'batteries-included' platform
 - We leverage opinionated Red Hat stack (e.g. prefer supported operators)
- We don't lift and shift. Modernize Apps first
- Migrating Legacy Microsoft HPC workloads this year

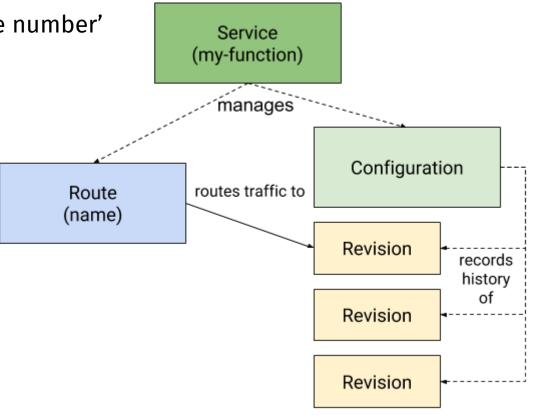


Openshift Serverless

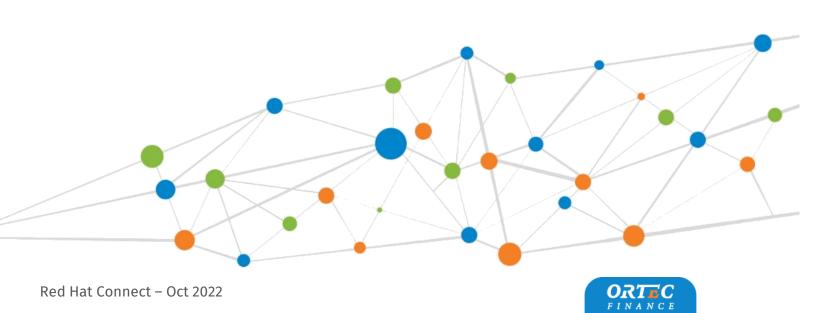
The ability to scale to zero – "Some people call this Serverless"

o Benchmark:

- dotnet core implementation of naive 'highest prime number'
- We kept it synchronous (no Knative Eventing)
- Node selectors enabled



DEMO









ARO VM: Standard_E4s_v3 ROSA VMs: m5.xlarge

Benchmark: hey -z 10s -c 30 "<app_url>/100000"

ARO slightly faster for same workload on different VM type

```
Summary:
 Total: 10.0386 secs
 Slowest: 3.1927 secs
 Fastest: 0.0088 secs
 Average: 0.0502 secs
 Requests/sec: 595.8994
 Total data: 173478 bytes
 Size/request: 29 bytes
Response time histogram:
 0.009 [1]
 0.327 [5951]
 0.646 [0]
 0.964 [0]
 1.282 [0]
 1.601 [0]
 1.919 [0]
 2.238 [0]
 2.556 [0]
 2.874 [0]
 3.193 [30]
```

```
Summary:
 Total: 10.0238 secs
 Slowest: 5.2999 secs
 Fastest: 0.0206 secs
 Average: 0.0581 secs
 Requests/sec: 515.2713
 Total data: 149791 bytes
 Size/request: 29 bytes
Response time histogram:
 0.021 [1]
 0.549 [5134]
 1.076 [0]
 1.604 [0]
 2.132 [0]
 2.660 [0]
 3.188 [0]
 3.716 [0]
 4.244 [0]
 4.772 [0]
 5.300 [30]
```

ARO vs Rosa performance review II

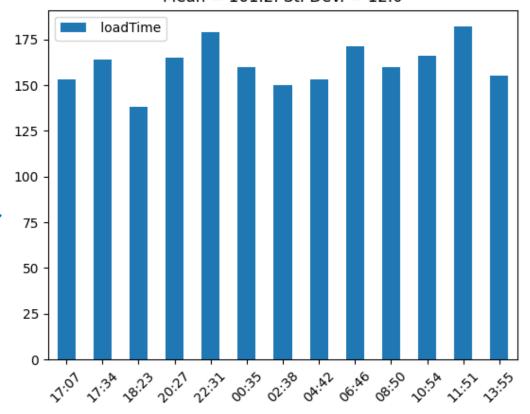
ROSA:

Provisioning time per machine family in seconds		
M5.xlarge	201 s	

ARO:

Provisioning time per mac	nine family in seconds	
Standard_D4s_v3	161 s	
Compute optimized 32 core	es	
(F32)	230 s	
Compute optimized 16 core	es	
(F16)	232s	
Compute optimized 8 core	s (F8) 244 s	

Mean = 161.2. St. Dev. = 12.0



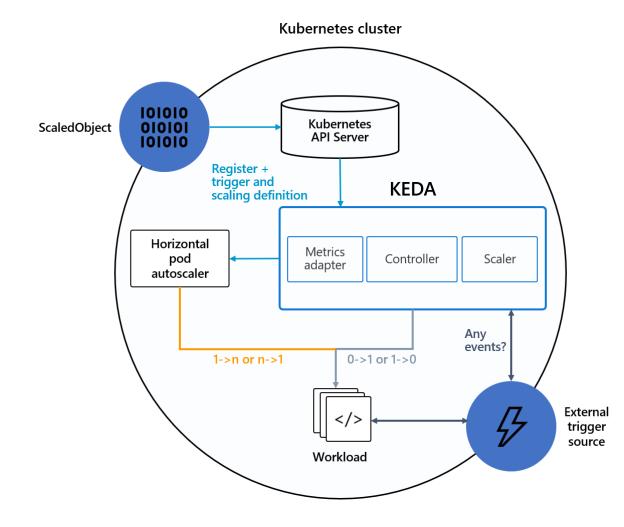


KEDA

Kubernetes Event Driven Autoscaler

o Benchmark:

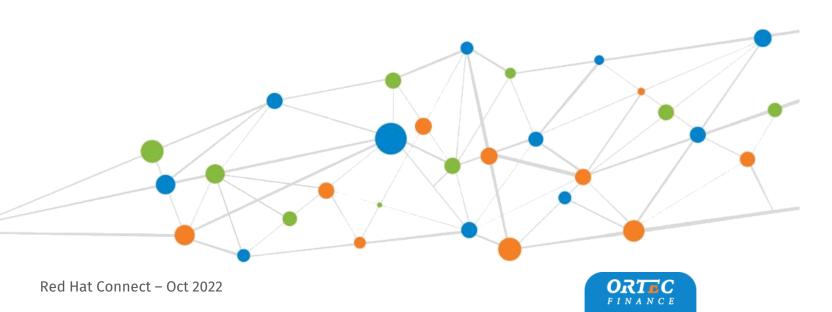
- Same prime number algo
- Async; Producer / Consumer set up
- No Kafka, but AMQ Broker







DEMO











- Operates on standard k8s resources
- Can scale existing deployed apps
- Pull based approach
- Doesn't manage data delivery
- K8s Horizontal Pod Autoscaler (HPA)
- Focus is on event driven autoscaling

- Operates on Knative Service
- Existing apps must be converted
- Push based approach
- Manages data delivery (Eventing)
- Knative Autoscaler
- Demand-based autoscaling (HTTP)





Wrap-up - Performance ROSA vs ARO

- ARO provisioning performance is likely to be ~10% better.
- ROSA supports Spot instances for machinesets
- ROSA machinesets can only be configured in the OCM or with the rosa cli
- ROSA cluster can hold up to 180 nodes while ARO cluster only scale to 40
- Default worker machine types are very similar in performance for both ROSA and ARO
- Long running workloads can rely on Machine autoscaling (ROSA had troubles scaling up from 0 machines)
- Short running workloads will rely on 'one hot spare'-approach



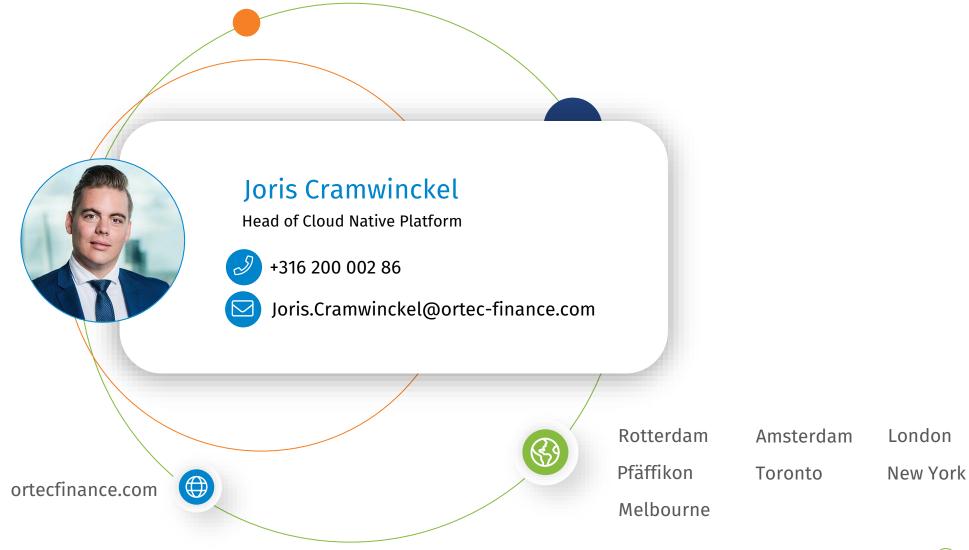


Wrap-up - Scaling Pods from 0 to 100s

- o For (long) parallel multiproces computations we prefer the **KEDA** setup
 - + Operates on standard K8 resources
 - + Can scale existing apps
 - + Tweakable autoscaling (e.g. deal with cold starts)
 - Not all community KEDA scalers not production ready (EDA for Prometheus is shipped in 4.11)
- For (short) multithreaded Compute we want to investigate **Openshift Serverless** further:
 - + Simplified deployment syntax
 - + Also includes traffic distribution
 - Async support requires rigorous design
 - Not designed for long running tasks (e.g. Kourier time-outs in 300s)
 - Overshoots in #pods (in default settings)



Contact me



Red Hat Connect – Oct 2022

Public

Public